



AUBURN

UNIVERSITY

SAMUEL GINN
COLLEGE OF ENGINEERING

Binary Heaps

Motivation: Priority Queue

Conceptually similar to a stack or a queue.

A **priority queue** chooses the next element to delete based on **priority**.



The element returned by the remove operation will be the one with the most **extreme priority** (max or min, depending on how the priority queue is configured).

Priority is some value associated with the element that could represent importance, cost, or some other problem-specific concept.

Applications:

Interrupt handling, bandwidth management, simulation, sorting, graph algorithms, selection algorithms, compression algorithms

Priority Queue

PQ Method	Unsorted List	Sorted List	Balanced BST
add	$O(1)$	$O(N)$	
remove	$O(N)$	$O(1)$	
peek	$O(N)$	$O(1)$	





Priority Queue

PQ Method	Unsorted List	Sorted List	Balanced BST
add	$O(1)$	$O(N)$	} ?
remove	$O(N)$	$O(1)$	
peek	$O(N)$	$O(1)$	

Q. What is the worst-case for each PQ if using a balanced BST?

	A	B	C	D
add	$O(N)$	$O(N)$	$O(\log N)$	$O(\log N)$
remove	$O(N)$	$O(1)$	$O(\log N)$	$O(\log N)$
peek	$O(N)$	$O(1)$	$O(1)$	$O(\log N)$

Priority Queue

PQ Method	Unsorted List	Sorted List	Balanced BST	Binary Heap
add	$O(1)$	$O(N)$	$O(\log N)$	$O(\log N)$
remove	$O(N)$	$O(1)$	$O(\log N)$	$O(\log N)$
peek	$O(N)$	$O(1)$	$O(\log N)$	$O(1)$
				
	<i>Nodes or arrays</i>	<i>Nodes or arrays</i>	<i>AVL, R-B, etc.</i>	<i>Nodes or arrays</i>

← But...

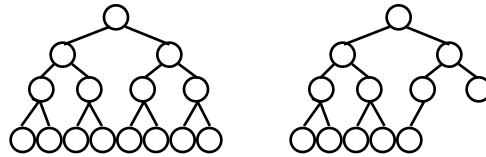
An array-based implementation is the most common and preferred.

The term “heap” usually implies an array.

Binary heaps

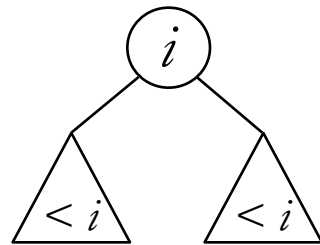
Binary heaps

A binary heap is a **complete binary tree** ...



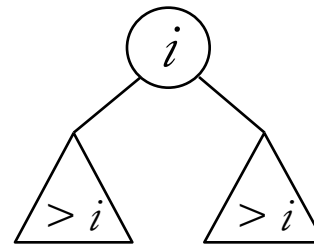
Height is $O(\log N)$

... in which each node obeys a **partial order property**.

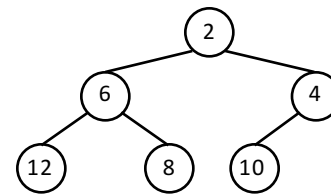
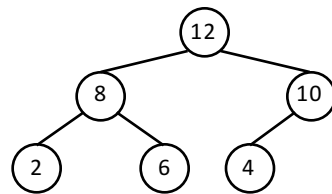


max heap

or



min heap



Array-based implementation

Binary heaps are usually implemented as an array because:

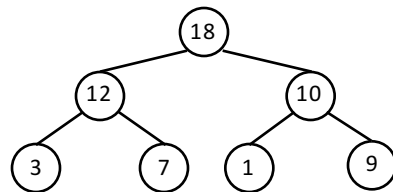
Acceptably space efficient (complete shape).

Easy traversal: parent to child via multiplication, child to parent via division

Many ways to map a hierarchy onto a linear array, but this is the one that we will use:

- Store the root at index 0.
- For a node stored at index i , store its left child at index $2i+1$ and its right child at index $2i+2$. Thus, the parent of a node stored at index i will be at index $(i-1)/2$.

Conceptually:



Implemented:

18	12	10	3	7	1	9
0	1	2	3	4	5	6

Adding values

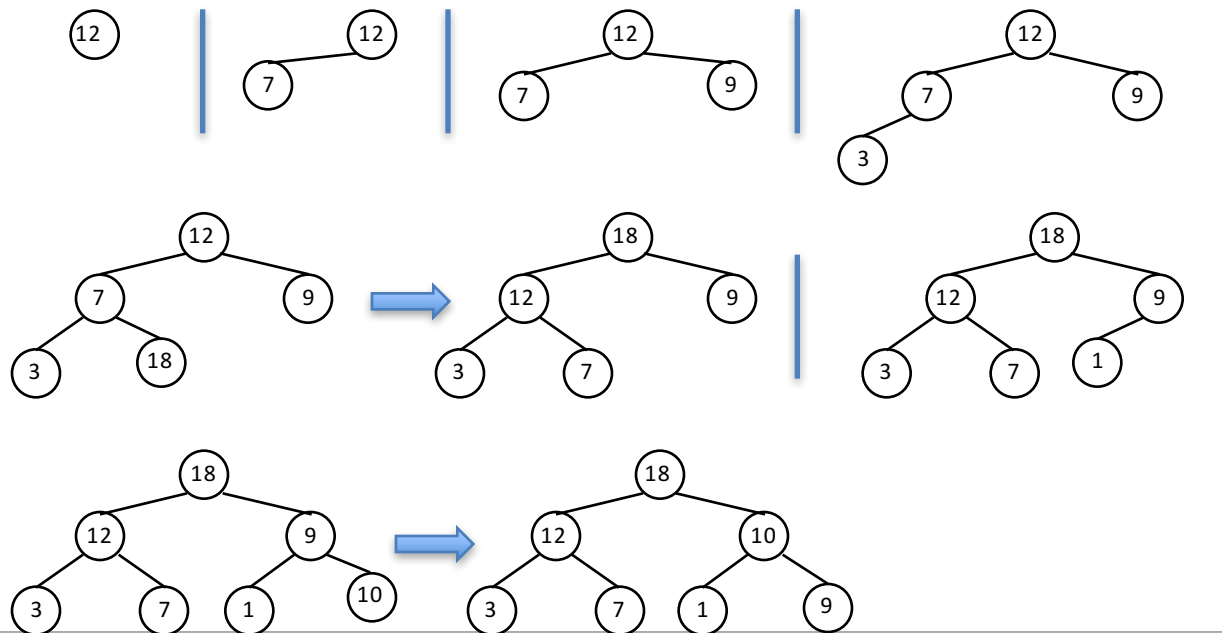
1. Insert the new element in the one and only one location that will maintain the complete shape.
2. Swap values as necessary on a leaf-to-root path to maintain the partial order.

Max heap example: 12, 7, 9, 3, 18, 1, 10

Adding values

1. Insert the new element in the one and only one location that will maintain the complete shape.
2. Swap values as necessary on a leaf-to-root path to maintain the partial order.

Max heap example: 12, 7, 9, 3, 18, 1, 10



Adding values

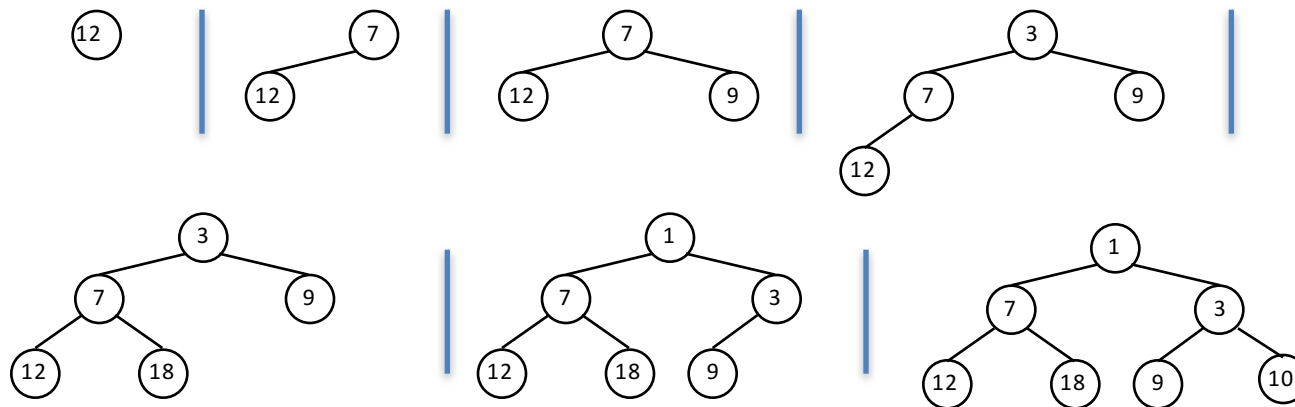
1. Insert the new element in the one and only one location that will maintain the complete shape.
2. Swap values as necessary on a leaf-to-root path to maintain the partial order.

Min heap example: 12, 7, 9, 3, 18, 1, 10

Adding values

1. Insert the new element in the one and only one location that will maintain the complete shape.
2. Swap values as necessary on a leaf-to-root path to maintain the partial order.

Min heap example: 12, 7, 9, 3, 18, 1, 10

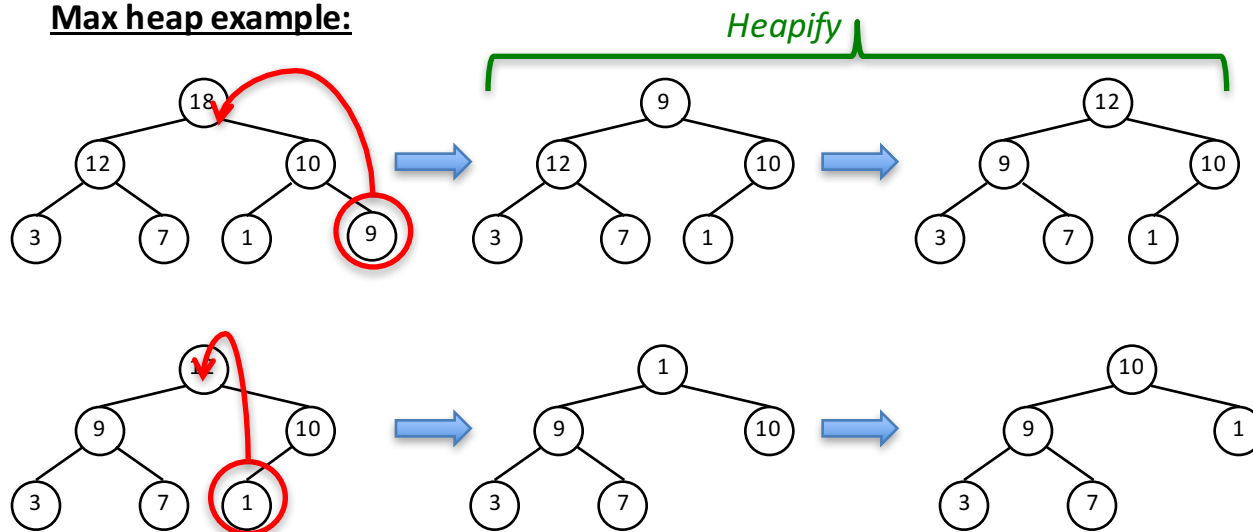


Removing values

Delete and return the element with the extreme (max/min) priority.

1. Maintain the complete shape by replacing the root value with the value in the lowest, right-most leaf. Then delete that leaf.
2. Swap values as necessary on a root-to-leaf path to maintain the partial order.

Max heap example:



Application: sorting

Heapsort

Heapsort works in two phases: (1) The initial arbitrary order of the array is transformed into a partial order, and then (2) the partial order is transformed into a total order.

1. Rearrange the array elements into max heap order.
2. Repeatedly move the maximum element to its final sorted place toward the end of the array, and heapify the remaining elements.

Example:

20	12	35	15	10	80	30	17	2	1
0	1	2	3	4	5	6	7	8	9



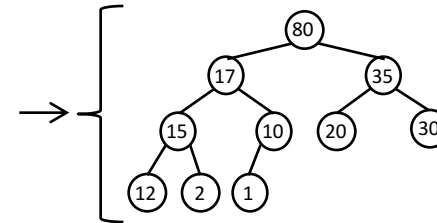
$O(N \log N)$ *Actually, $O(N)$*

80	17	35	15	10	20	30	12	2	1
0	1	2	3	4	5	6	7	8	9



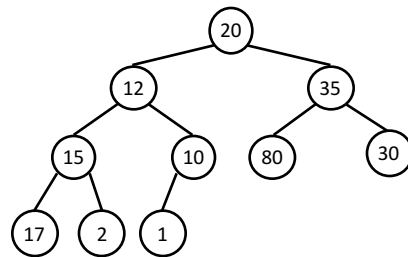
$O(N \log N)$

1	2	10	12	15	17	20	30	35	80
0	1	2	3	4	5	6	7	8	9



Heapsort

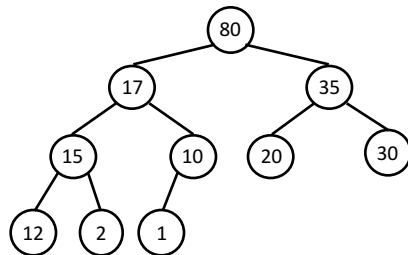
20	12	35	15	10	80	30	17	2	1
0	1	2	3	4	5	6	7	8	9



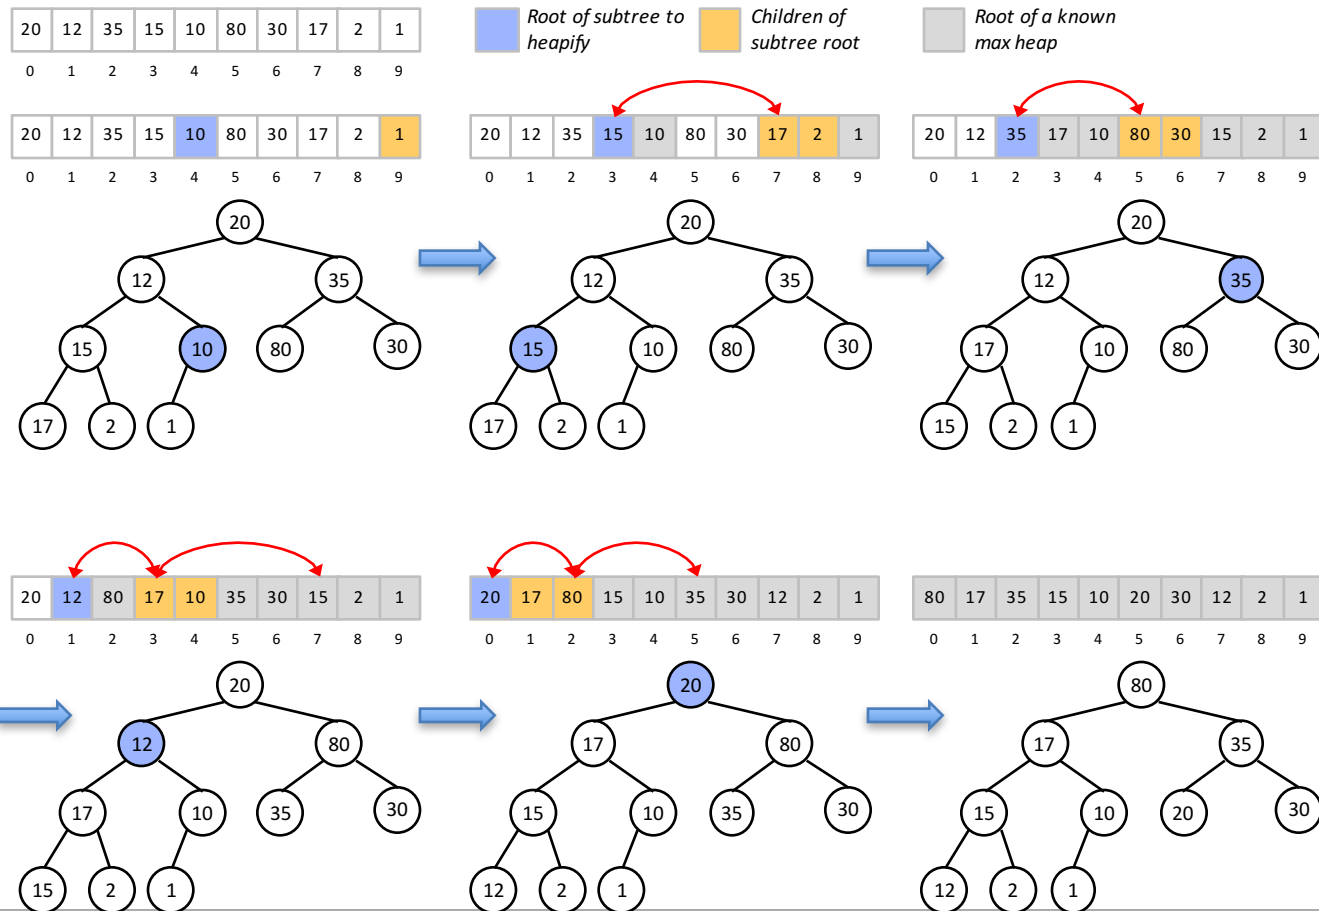
1. Rearrange the array elements into max heap order.

Beginning with the lowest, right-most parent and continuing to the root, heapify each subtree.

80	17	35	15	10	20	30	12	2	1
0	1	2	3	4	5	6	7	8	9

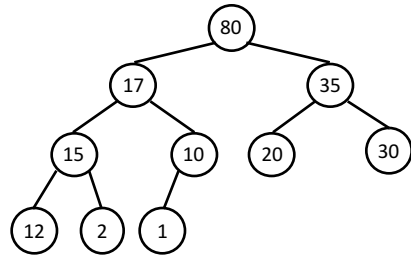


Heapsort



Heapsort

80	17	35	15	10	20	30	12	2	1
0	1	2	3	4	5	6	7	8	9

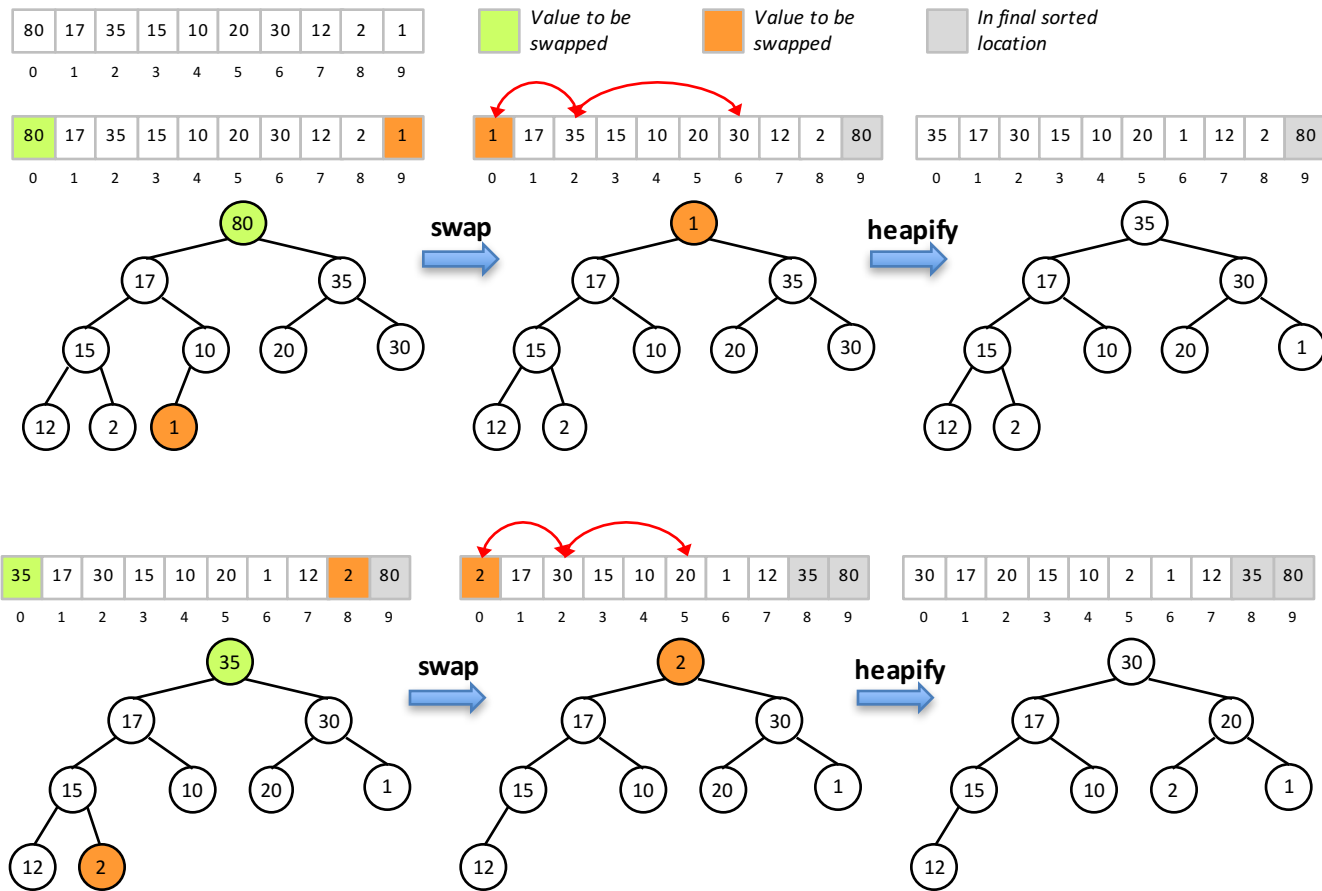


2. Repeatedly move the maximum element to its final sorted place toward the end of the array, and heapify the remaining elements.

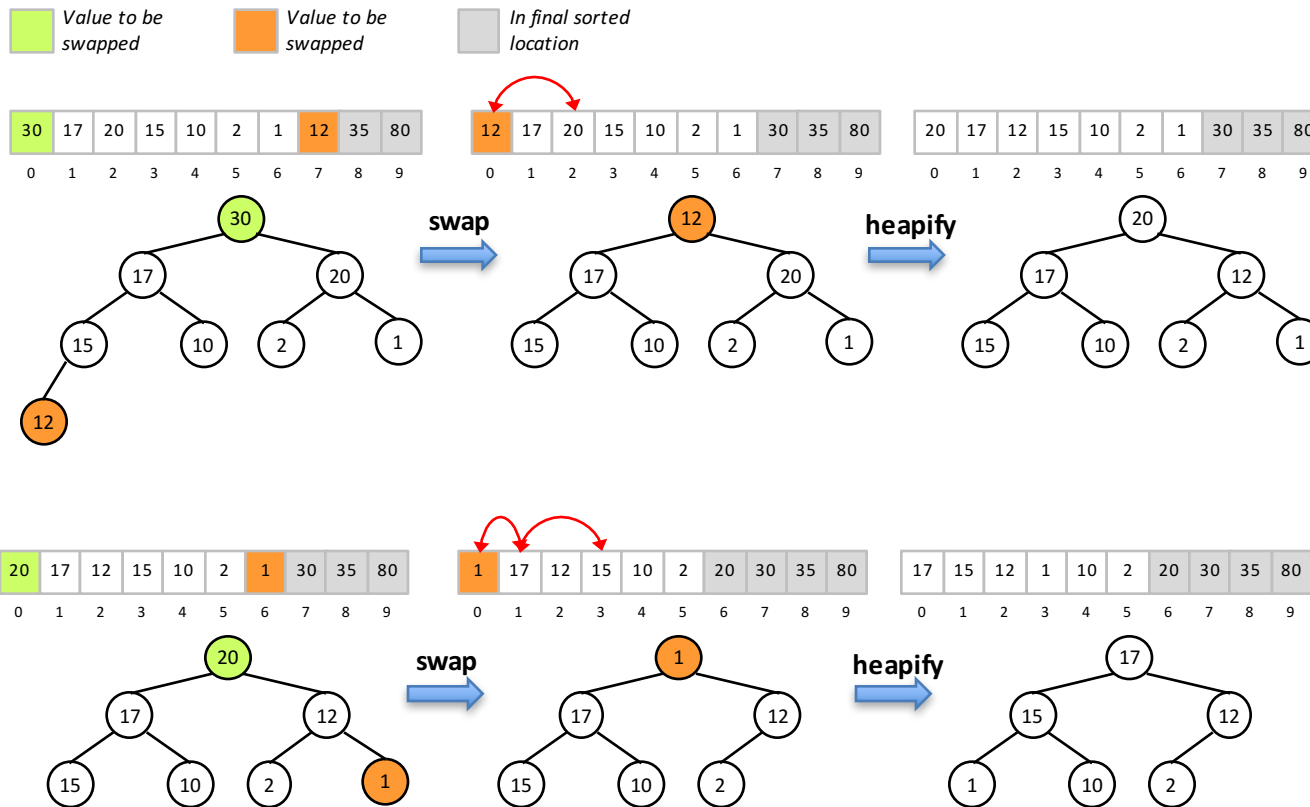
Set last to a.length - 1
Swap a[0] and a[last]
last--
Heapify a[0 .. last]
Repeat until last == 0

1	2	10	12	15	17	20	30	35	80
0	1	2	3	4	5	6	7	8	9

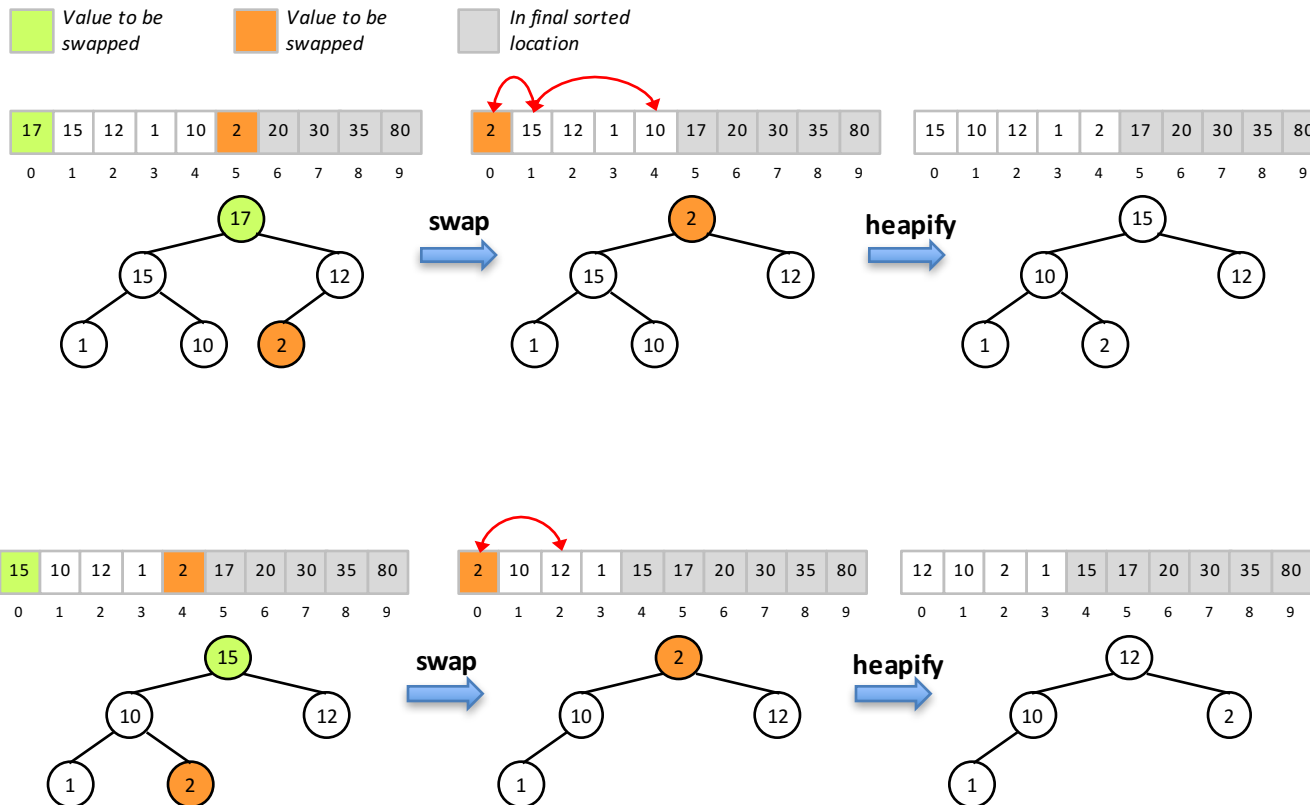
Heapsort



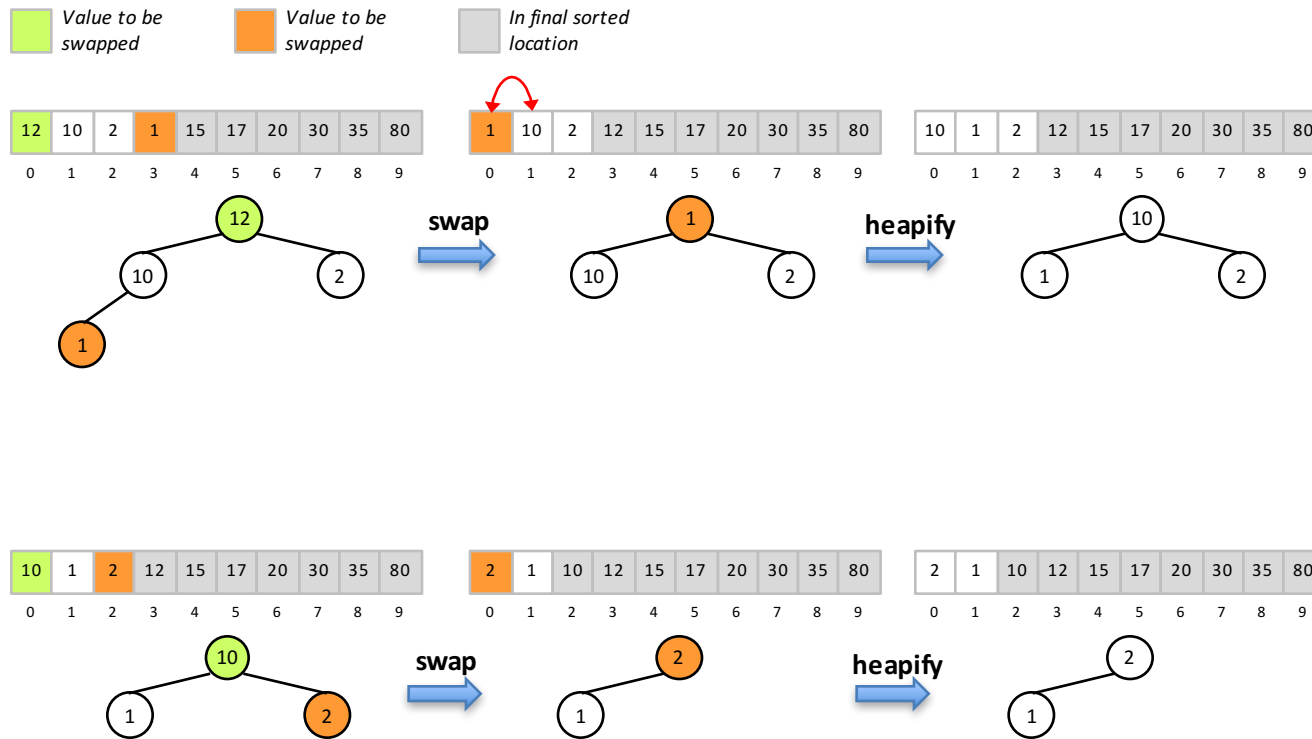
Heapsort



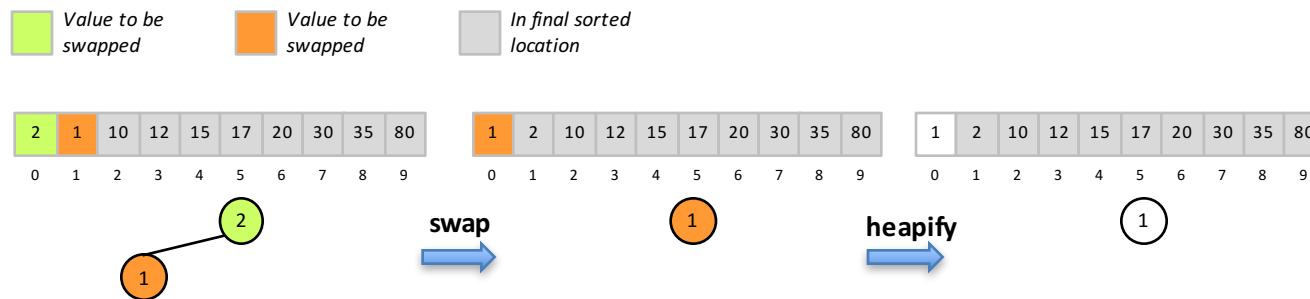
Heapsort



Heapsort



Heapsort



Heapsort is an in-place sort with guaranteed $N \log N$ worst-case performance.

Mergesort? No. ($N \log N$ worst-case, but needs N extra space.)

Quicksort? No. (N^2 in worst-case.)

But, heapsort is not stable and it typically has larger constant factors than quicksort.

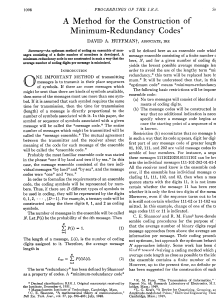
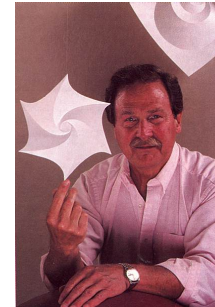
Application: Huffman's algorithm

Huffman's algorithm

Huffman's algorithm generates a variable-length encoding for a given alphabet for the purposes of data compression.

Developed by [David Huffman](#) in 1951 as a class project at MIT, and published in 1952.

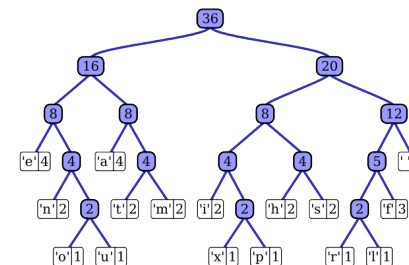
Widely used today as part of various compression utilities (PKZIP, MP3, JPEG).



Famous story:

In 1951, David A. Huffman and his MIT information theory classmates were given the choice of a term paper or a final exam. The professor, Robert M. Fano, assigned a term paper on the problem of finding the most efficient binary code. Huffman, unable to prove any codes were the most efficient, was about to give up and start studying for the final when he hit upon the idea of using a frequency-sorted binary tree and quickly proved this method the most efficient.

In doing so, the student outdid his professor, who had worked with information theory inventor Claude Shannon to develop a similar code. Huffman avoided the major flaw of the suboptimal Shannon-Fano coding by building the tree from the bottom up instead of from the top down.



Huffman's algorithm

ASCII American Standard Code for Information Interchange

Binary character encoding scheme: A sequence of 0s and 1s (bits) used to encode characters.

ASCII includes English alphabet, punctuation, digits, and “control” characters (e.g., newline, carriage return).

The 95 printable characters in ASCII:

!"#\$%&'()*+,-./0123456789:;<=>?@ABCDEFGHIJKLMNOPQRSTUVWXYZ[\]^_`abcdefghijklmnopqrstuvwxyz{|}~

Binary	Decimal	Char
100 0000	64	@
100 0001	65	A
100 0010	66	B
100 0011	67	C
100 0100	68	D

ASCII is a fixed length code. Each character is represented by the same number of bits.

In US-ASCII, each character is represented in one byte (8 bits).

```
% more abcfile.txt
ABC
% ls -l abcfile.txt
-rw-r--r--  1 User  User   4 Apr  2 09:18 abcfile.txt
%
```

8 bits = 1 parity bit and 7 bits to encode the character. $2^7 = 128$ different characters

Huffman's algorithm

Text file:

BABACEDABCDABABACD
ADABCCABCA

28 characters

Text compression
stores the same
information in fewer
bytes.

Binary ASCII form:

01000010010000010100
00100100000101000011
01000101010001000100
00010100001001000011
01000100010000010100
0010...

28 bytes

$28 * 8 = 224$ bits

We could compress
this file by taking
advantage of the fact
that some characters
appear more often
than others.

Huffman's algorithm

Number of bits per character determined by the char's relative frequency of occurrence.

Most frequently occurring characters should use the fewest bits.

Text file:

BABACEDABCDABABACD
ADABCCABCA

Character frequency:

A-10, B-7, C-6, D-4, E-1

A variable length code:

A = 11
B = 10
C = 00
D = 011
E = 010

"Compressed" file:

10111011000100111110
00011111011101100011
11011111000001110001
1

Only 61 bits

Uncompressed file required 196 bits

**This would compress the file to 31%
of its original size.**

Huffman's algorithm

A first attempt: Iterate over the alphabet in descending order of frequency. Assign the next smallest unique bit string to the current character, starting with '0'.

Character frequency:

A-10, B-7, C-6, D-4, E-1

The variable length code:

A = 0

B = 1

C = 01

D = 10

E = 11

The vlc must have the
prefix property.

The code for one character can't be a prefix of another character's code.

Text file:

BABACEDABCDABABACD

ADABCCABCA

↓ *zip*

"Compressed" file:

1010011110010110...

↓ *unzip* *Can't reconstruct the original!*

Does the file start with a B or a D??

Huffman's algorithm

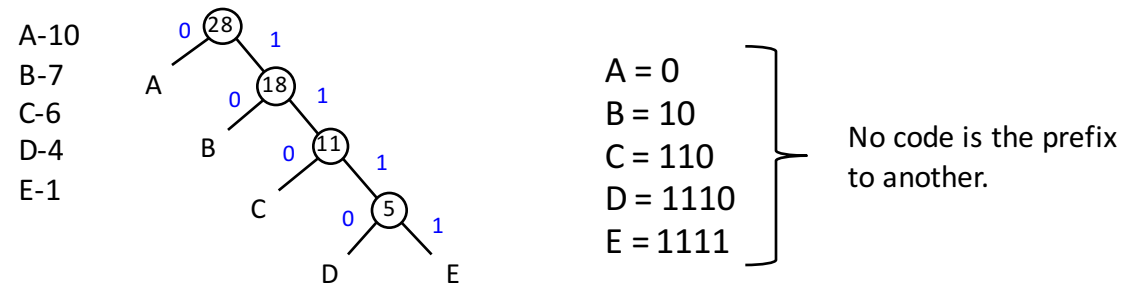
Binary trees in which the leaves contain the characters to be coded.

Interior nodes are just place-holders.

The root of every subtree is annotated with the cumulative frequency of all its descendent leaves.

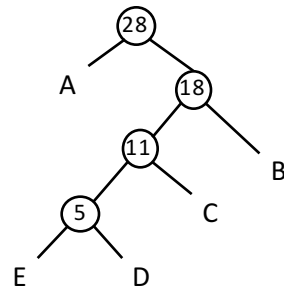
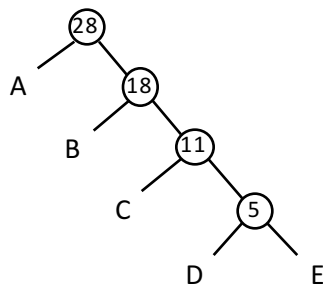
Character codes are generated by root to leaf traversals.

Left branch = 0, Right branch = 1



Huffman's algorithm

A-10, B-7, C-6, D-4, E-1



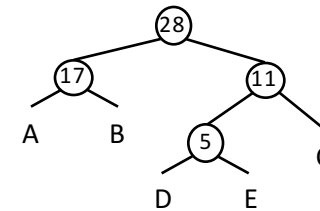
• • •

We would like to use the code tree with minimum **expected code length**.

$$L(C) = \sum_{i=1}^N w_i \times \text{length}(c_i)$$

This is just a weighted average of all possible character code lengths.

A:	$(10 \div 28) * 1 = 0.36$
B:	$(7 \div 28) * 2 = 0.50$
C:	$(6 \div 28) * 3 = 0.66$
D:	$(4 \div 28) * 4 = 0.56$
E:	$(1 \div 28) * 4 = 0.12$
	<u>2.20</u>



2.18

Huffman's algorithm generates a code tree with an expected code length that is at least as small as any other code tree that could be generated.

Huffman's algorithm

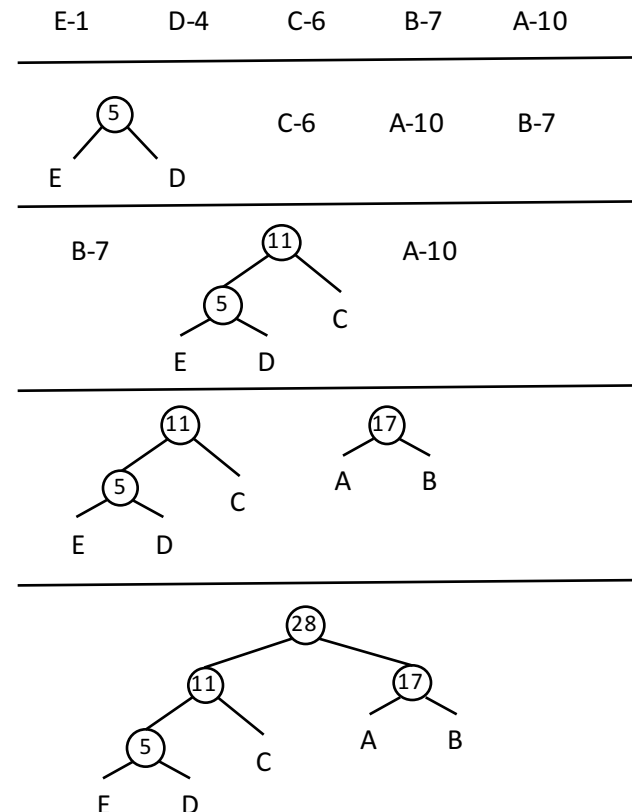
Generates a variable length code with the prefix property such that there is no other encoding with a smaller expected code length.

A-10, B-7, C-6, D-4, E-1

Create a single node code tree for each character and insert each of these trees into a priority queue (min heap).

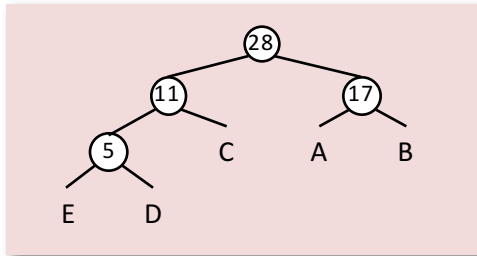
```
while (pq has more than one element) {
    c1 = pq.deletemin();
    c2 = pq.deletemin();
    c3 = new codetree(c1,c2);
    pq.add(c3);
}
```

Char	Encoding
A	10
B	11
C	01
D	001
E	000



Huffman's algorithm

A-10, B-7, C-6, D-4, E-1

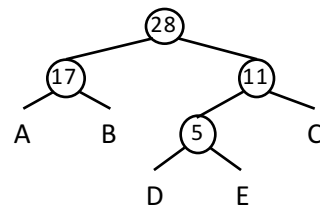


Char	Encoding
A	10
B	11
C	01
D	001
E	000

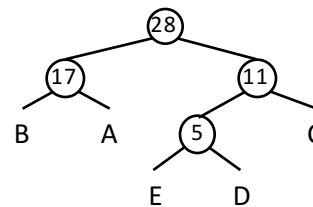
Expected code length:

$$\begin{aligned}
 A: & (10 \div 28) * 2 = 0.71 \\
 B: & (7 \div 28) * 2 = 0.50 \\
 C: & (6 \div 28) * 2 = 0.43 \\
 D: & (4 \div 28) * 3 = 0.43 \\
 E: & (1 \div 28) * 3 = 0.11 \\
 & \underline{\hspace{1.5cm}} \\
 & 2.18
 \end{aligned}$$

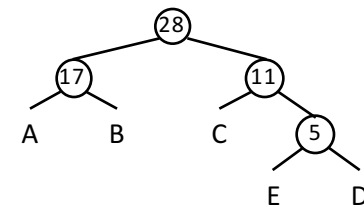
This is not the only code tree with minimum expected code length.



Char	Encoding
A	00
B	01
C	11
D	100
E	101



Char	Encoding
A	01
B	00
C	11
D	101
E	100

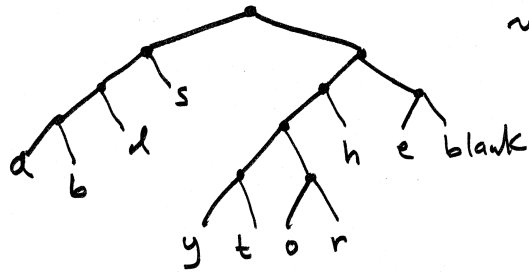


Char	Encoding
A	00
B	01
C	10
D	110
E	111

Huffman's algorithm

She sells sea shells by the sea shore

~44%



I slit the sheet, the sheet I slit, and on the slitted sheet I sit

~46%

